

# UNIT - 1 | Transaction Management

**Transaction management** ensures data integrity in systems that involve changes to multiple data elements. It's a process that oversees database operations to guarantee data consistency. It ensures all parts of a data change are completed successfully or none are applied.

**Transaction:** A single unit of work in a database system. Represents a series of database operations (reads, writes, etc.).

**Savepoints:** Mark a point within a transaction where a rollback can occur.

## Transaction States

**Transaction States** describe the different stages a transaction goes through in its lifecycle within a Database Management System. It ensures data consistency during database operations.

- **Active State:** Executing instructions (reads, writes).
- **Partially Committed State:** Changes temporary (not yet on disk).
- **Committed State:** Successful completion (changes permanent).
- **Failed State (Aborted State):** Error occurred (changes rolled back).

## ACID Properties

ACID ensures data reliability in databases. It treats multiple operations as one, guaranteeing a consistent state even in case of failures.

- **Atomicity:** Transactions are all-or-nothing. Everything happens or nothing does.
- **Consistency:** Transactions move the database from one valid state to another.
- **Isolation:** Concurrent transactions don't interfere with each other.
- **Durability:** Committed transactions survive system failures.

## Operations in Transactions

Transactions in a DBMS are a series of logically related operations that ensure data integrity.

### Data Operations:

- **Read:** Grab data (e.g., checking account balance).
- **Write:** Modify data (e.g., updating balance after transfer).

### Transaction Control:

- **Commit:** Finalize changes (e.g., permanent transfer).
- **Rollback:** Undo changes (e.g., insufficient funds during transfer).

## Storage Structures

DBMS stores data in specific formats that affect how quickly we can access and use it.

### Types of Storage Structures:

**1. File Organization:** How data is physically stored in a file.

- **Heap Files:** Fast for adding data, slow for searching.
- **Sorted Files:** Fast for searching by key, slow for inserts (requires re-sorting).
- **Hashed Files:** Fast for retrieving specific data by key, less efficient for finding ranges.

**2. Indexed Organization:** Improves access speed using indexes.

- **B-Trees:** Efficient for searching ranges and specific data.
- **ISAM:** Good for both searching by key and sequential access.

## Concurrent Executions

The ability of a DBMS to process multiple transactions simultaneously in a shared database.

**Benefits:** Increased Throughput, Reduced Wait Time, Improved Resource Utilization.

Concurrent execution can lead to data inconsistencies if not managed properly.

**Interleaved Execution:** Transactions might access and modify the same data items at the same time, leading to unpredictable results.

## Serializability

Serializability ensures that concurrent transactions on a database appear to run one after another, even if they happen at the same time.

**1. Conflict Serializability:** This is stricter and avoids data inconsistencies by ensuring conflicting operations on the same data item happen in a specific order (read before write, or write before read).

**2. View Serializability:** This is more relaxed and focuses on the final outcome. As long as the final state of the database is consistent, the exact order of conflicting operations within a transaction doesn't matter. This allows for more concurrent access.

## Concurrency Control

Concurrency control is a mechanism in DBMS that manages concurrent access to data by multiple users or applications. It ensures data consistency and integrity in a multi-user environment.

**Goals of Concurrency Control:** ACID properties, Serializability

**Concurrency Problems:** When multiple users edit the same data at the same time in a database, things can get messy (data inconsistencies).

**Types of Concurrency Problems:**

- **Lost Updates:** Two updates collide, and only one wins.
- **Inconsistent Reads:** Reading data twice shows different results due to mid-air updates.
- **Phantom Reads:** New or missing data appears/disappears between reads due to inserts/deletes by others.
- **Dirty Reads:** Reading data that might be later discarded (Seeing Unfinished Work).

**Preventing Concurrency Problems:**

- **Locking Mechanisms:** Locks prevent conflicts during data updates.
- **Timestamp Ordering:** Timestamps ensure a specific execution order.
- **Optimistic Concurrency Control:** Allows concurrency but checks for conflicts later.

- **Database Isolation Levels:** Define the degree of isolation between transactions.

## Concurrency Control Protocols

Mechanisms to manage concurrent access to data by multiple users or processes. Their goal is to ensure data consistency and integrity despite these simultaneous operations.

### 1. Lock-Based Protocols

- Transactions acquire locks on specific data items (records, tables) before modifying them.
- Prevents other transactions from conflicting modifications.

#### Modes of Locks:

- **Exclusive Lock (X-lock):** The exclusive access holding transaction can modify the data.
- **Shared Lock (S-lock):** Many users can read, but no one can write.

**Granting of Locks:** The lock manager determines if a lock request can be granted based on the current locking state and the type of lock requested (X or S).

- **Granting an X-lock:** Only granted if no one else has any lock (read or write) on the data.
- **Granting an S-lock:** Okay if others already have S-locks (read-only access), but blocks any exclusive lock requests.

**Two-Phase Locking Protocol (2PL):** 2PL ensures a transaction's locking actions have a well-defined order, preventing inconsistencies.

- **Growing Phase (Locking Up):** Grab locks on needed data. No releasing locks yet.
- **Shrinking Phase (Releasing Locks):** Release unused locks. No grabbing new ones now.

**Lock Point:** The Point at which the growing phase ends.

### 2. Timestamp-Based Protocols

- Each transaction receives a unique timestamp when it starts.
- Transactions are serialized based on their timestamps, ensuring a specific execution order.

#### Timestamp Ordering Protocol:

- Assigns a unique timestamp (TS) to each transaction.
- Ensures serializability by ordering transactions based on their timestamps.

#### Thomas Write Rule (Modification of Timestamp Ordering Protocol):

- A concurrency control mechanism used in database systems to ensure data consistency.
- Improves efficiency over basic timestamp ordering.
- Allows ignoring outdated writes, improving performance compared to strictly rolling back transactions

### 3. Validation-Based Protocols

- Transactions are validated after their execution, checking for conflicts with previously committed transactions. If a conflict is found, the transaction is rolled back (undone).

#### Process:

- **Read Phase:** Transaction reads data and stores it in local copies.
- **Validation Phase:** Local copies are compared against actual data for conflicts.
- **Write Phase (if validation succeeds):** Local copies update the actual database.

## Deadlock Handling

Deadlocks occur in a DBMS when two or more transactions wait indefinitely for resources (locks) held by each other. This creates a circular dependency, halting all involved transactions.

- 1. Deadlock Prevention:** DBMS analyzes transactions to ensure resource allocation never leads to a deadlock scenario. Prevents deadlocks altogether.
- 2. Deadlock Avoidance:** DBMS maintains information about resource requests and uses algorithms to predict and avoid deadlocks. Prevents potential deadlocks during transactions.
- 3. Deadlock Detection and Recovery:** DBMS monitors transactions and identifies wait-for graphs to detect deadlocks. Allows deadlocks to occur but detects them for resolution.

## UNIT - 2 | Database System Architectures

Database architecture is the blueprint for how data is stored, organized, accessed, and managed. It's like the foundation of a house, crucial for smooth data operations.

### Types of Database Architectures:

#### 1. Tier Architecture:

- **1-Tier (Single-tier):** Easiest to develop, but least scalable.
- **2-Tier:** Separates user interface from data storage. More scalable.
- **3-Tier:** Most scalable, with a dedicated layer for handling complex business logic

#### 2. Data Model:

- **Relational Model:** Data is stored in tables with rows and columns. Relationships between tables are established through foreign keys.
- **NoSQL Model:** More flexible, suitable for large unstructured datasets.

## Centralized Architecture

All data and processing logic reside on a single server. Clients (usually dumb terminals) connect and interact directly with the server.

**Pros:** Simple to set up and manage. Cost-effective for small databases.

**Cons:** Limited scalability, Single point of failure.

## Client-Server Architecture

Data is stored on a central server, but clients (PCs, workstations) handle user interface and some processing. Clients communicate with the server to access and manipulate data.

**Pros:** More scalable, Improved performance, More flexible.

**Cons:** More complex to set up and manage, Requires robust network infrastructure.

# Server System Architecture

Refers to the organization of processes and data storage on the server side.

- **Transaction Servers:** Used in relational databases, these handle queries and transactions efficiently.
- **Data Servers:** Used in object-oriented databases, these focus on data access and processing on powerful client machines.

## 1. Transaction Server Process Structure:

Consists of multiple processes that manage user requests (transactions), data access, and overall system functionality.

### Components:

- **Server Processes:** Handle user requests, run them concurrently, and return results.
- **Lock Manager Process:** Ensures data consistency by controlling access.
- **Database Writer Process:** Periodically saves changes from memory to disk.
- **Log Writer Process:** Tracks all database modifications for recovery.
- **Checkpoint Process:** Creates regular backups for faster disaster recovery.
- **Process Monitor Process:** Stores frequently used data for faster access.
- **Shared Memory:** Keeps an eye on everything and restarts processes if needed.

## 2. Data Servers:

Prioritize data access & processing on client machines. Often used in object oriented databases.

### Key Points:

- Clients have more processing power and handle complex queries locally.
- Data servers can ship data (pages or entire objects) to clients for processing.
- Data caching on clients improves performance for frequently accessed data.

### Key Differences:

- **Focus:** Transaction servers handle user requests and data consistency, while data servers manage physical storage and access.
- **Location:** Transaction servers are often close to clients, while data servers can be separate for better scalability.

# Parallel Systems

Parallel DBMS utilizes multiple processors and disks to improve database performance.

### Benefits:

- **Increased Speed:** Queries and operations run concurrently, leading to faster results.
- **Improved Scalability:** allows the system to handle larger datasets & more users efficiently.
- **Enhanced Availability:** If one node fails, others can continue processing.

### Architectural Designs:

- **Shared Memory Architecture:** Processors share global memory space for efficient access.

- **Shared Disk Architecture:** Processors have private memories but access shared storage devices for data.
- **Shared-Nothing Architecture:** Each node has its own CPU, memory, and storage.

## Speed up parallel systems

- Focuses on **reducing execution time** for a fixed workload.
- Achieved by adding more processing resources (CPUs, machines) to a system.
- Ideally, speedup is linear with the number of resources added.

## Scale up parallel systems

- Focuses on handling an **increasing workload** while maintaining performance.
- Achieved by adding resources proportionally to the workload growth.
- Ensures the system can accommodate larger datasets and more users without performance degradation.

## Interconnection Networks

Defines how processors or storage devices communicate in a Database Management System with parallel processing capabilities

**Bus:** Simple shared channel for all devices. Only one device transmits at a time.

**Mesh:** Processors arranged in a grid, connected to nearest neighbors. Offers multiple paths for data transfer, improving scalability.

**Hypercube:** Highly scalable network based on n-dimensional cubes. Processors connected only to nodes differing in a single dimension.

## Parallel Database Architectures

Parallel databases distribute tasks across multiple processors to improve performance.

1. **Shared Memory Architecture:** Tightly coupled system with all processors sharing a single memory space. Data and code reside in the shared memory, accessible by all processors.
2. **Shared Disk Architecture:** Multiple processors connected to a shared storage device (disk). Data resides on the shared disk, while processors have local memory.
3. **Shared-Nothing Architecture:** Independent nodes (CPU, memory, storage) handle data subsets. Scales well, but complex for multi-location queries.
4. **Hierarchical Architecture:** Mix of shared-memory and shared-nothing. Upper level acts as a coordinator, distributing tasks to lower-level nodes.

## UNIT - 3 | Distributed Databases

A database system that stores data across multiple computers or sites on a network.

### Benefits:

- **Scalability:** Easily add more servers to handle growing data volumes or user access.
- **Availability:** If one site fails, others can still function (higher uptime).

- **Performance:** Distribute workload across multiple servers for faster queries.
- **Geographical Distribution:** Data can be located closer to users in different regions.
- **Improved fault tolerance:** Replication mechanisms ensure data consistency even in failure.

#### Components:

- **Sites:** Individual computers or servers storing database fragments.
- **Nodes:** Processing units within a site that manage data.
- **Distributed Database Management System (DDBMS):** Software that manages the entire distributed database and provides a unified view to users.

#### Distribution Schemes:

- **Fragmentation:** Dividing the database logically into smaller units for storage across sites.

#### Challenges:

- **Data Consistency:** Maintaining consistency of data across all sites after updates.
- **Complexity:** Managing data across multiple locations requires more complex software.
- **Network Reliability:** Reliant on a reliable network for communication between sites.

## Distributed Data Storage

Distributed data storage is a method of storing information across multiple physical servers, often spread across different locations. It breaks down data into chunks and distributes them for redundancy and improved performance.

**Benefits:** Scalability, Reliability, Performance, Availability (Similar to Distributed Databases).

#### Types of Distributed Storage:

- **Distributed File system (DFS):** Makes scattered data appear as a single drive (e.g., HDFS).
- **Distributed Block storage:** Splits data for fast transfers (e.g., SAN).
- **Distributed Object storage:** Stores data in self-contained units (e.g., Amazon S3).

**Applications:** Cloud storage services, Backup and disaster recovery etc.

## Distributed Transactions

A set of database operations executed across multiple interconnected databases (nodes) as a single logical unit. Ensures ACID properties even across distributed data.

#### Process (often uses Two-Phase Commit):

1. **Global Coordinator:** Initiates the transaction and coordinates participating nodes.
2. **Preparation Phase:** Each node involved prepares to commit changes.
3. **Commit Phase:** Coordinator instructs all nodes to commit.

If a node fails during the process, the transaction is aborted/rolled back.

#### Challenges:

- **Network failures:** Can disrupt communication and require complex rollback procedures.
- **Data consistency:** Ensuring all nodes reflect the same changes after a transaction.



# Commit Protocols

Commit protocols ensure data integrity and consistency in DBMS by coordinating transactions across multiple sites. They guarantee that a transaction either fully succeeds (commits) on all participating sites, or completely fails (aborts) on all sites.

**1. One-phase (local):** Simple commit for single-site databases.

**2. Two-Phase Commit (2PC):** Common choice, involves coordinator asking participants if ready, then committing or aborting based on replies.

- **Phase 1:** Preparing to Commit
- **Phase 2:** Committing or Aborting

**3. Three-phase (3PC):** Adds extra phase for participants to confirm readiness before final commit/abort. (More complex)

- **Phase 1:** Pre-commit (Similar to 2PC)
- **Phase 2:** Waiting for Agreement
- **Phase 3:** Commit or Abort

## Concurrency Control in Distributed Databases

- It **Ensures data consistency** when multiple transactions access the same data concurrently across distributed databases.
- Maintains **ACID properties** even in a distributed environment.
- Achieves **serializability**, preventing conflicts.

### Single Lock-Manager Approach:

- **Centralized coordinator** manages all locks for the distributed database.
- Transactions requesting a lock send a message to the lock manager.

### Distributed Lock Manager

- **Locks are managed locally** at each database site.
- Transactions coordinate locking across sites using protocols.
- **Protocols:** Two-phase commit (2PC) & Timestamp ordering.

## Parallel Databases

A database system that leverages multiple processors (CPUs) and disks to improve performance through parallel processing of tasks.

**Benefits:** Increased Throughput, Faster Queries, Improved Scalability, High Availability.

### Architectures:

- **Shared-Memory Architectures:** CPUs share memory, good for small setups.
- **Shared-Disk Architectures:** CPUs have private memory, share disk storage (more complex).
- **Shared-Nothing Architectures:** Each node has its own CPU, memory, disk.



# I/O Parallelism

Speeds up data retrieval from disk by dividing relations (tables) across multiple disks.

## Process:

- Data is partitioned based on a chosen attribute.
- Each partition is stored on a separate disk.
- Queries can then access data from multiple disks simultaneously.
- Results from each partition are combined after processing.

## Implementation:

- Requires a DBMS that supports parallel processing.
- Data partitioning strategy needs to be defined (e.g., by hash function, round-robin).

# Inter-Query Parallelism

The ability of a DBMS to execute multiple queries from various applications simultaneously.

- **Focuses on throughput:** Aims to handle a higher volume of transactions.
- **Independent execution:** Each query runs separately but concurrently with others.

# Intra-Query Parallelism

A technique to speed up complex queries by breaking them down into smaller, independent tasks that can be executed concurrently on multiple processors.

## Implementation:

- **Decomposing the Query:** The DBMS identifies independent operations within the query.
- **Parallelization:** These operations are assigned to different processors for simultaneous execution.
- **Combining Results:** The DBMS gathers and combines the partial results from each processor to generate the final query output.

# Intraoperation Parallelism

Intraoperation parallelism focuses on parallelizing individual operations within a single SQL query. This approach significantly improves query performance for large datasets.

**1. Parallel Sort:** Parallel sort divides the data into smaller chunks and sorts them concurrently on multiple processors.

**2. Parallel Join:** Parallel join partitions the tables and performs the join operation on each partition concurrently.

**Benefits:** Significant performance improvement for complex queries involving large datasets.

# Interoperation Parallelism

Interoperation parallelism deals with executing different operations within a single query expression in parallel.

**1. Pipelined parallelism:** Overlaps the execution of subsequent operations in a query.

**2. Independent parallelism:** Run independent tasks within a query simultaneously, like filtering separate tables.

#### **Benefits of Interoperation Parallelism:**

- Significantly reduces overall query execution time by overlapping operations.
- Makes efficient use of multiple processors in a system.

## **UNIT - 4 | PL/SQL**

Procedural Language extensions to SQL. Combines the power of SQL with procedural programming features (control flow, loops, etc.)

#### **Key Features:**

- **Procedural constructs:** control flow statements (if-then-else, loops) for complex logic
- **Cursors:** for iterating through result sets of SQL statements
- **Error handling:** exceptions to manage errors during program execution
- **Integration with SQL:** execute SQL statements directly within PL/SQL code

## **Advantages of PL/SQL**

- **Combines SQL and Procedural Code:** Logic and manipulate data within the same program.
- **Improved Performance:** Compiled code reduces network traffic and improves speed.
- **Increased Productivity:** Enables modular and reusable code.
- **Portable and Scalable:** Works across Oracle environments and handles complex tasks.
- **Object-Oriented Features:** Promotes better code organization.
- **Error Handling and Security:** Manages errors and restricts unauthorized access.
- **Web Application Development:** Can be used to create dynamic web content.

## **PL/SQL Blocks**

- Encapsulate a set of PL/SQL statements to perform a specific task.

#### **Structure:**

- **DECLARE:** Define variables, cursors, exceptions, etc. (Optional)
- **BEGIN:** Marks the start of executable statements.
- **Executable Statements:** Code to perform operations.
- **END:** Marks the end of the block.

## **Character Set**

- Set of characters a database can understand and store. Determines valid characters in identifiers, literals, and comments.

## **Literals**

- Fixed values directly included in your code.
- **Types:** Numeric, Character, String literals, Boolean.

## **PL/SQL Data Types**

**1. Scalar Data Types:** Represent single values.

- **NUMBER:** Stores numeric values with precision and scale.
- **BOOLEAN:** Represents true or false
- **CHAR/VARCHAR2:** Stores fixed or variable length character strings/
- **DATE:** Stores date and time information.

**Subtypes:** Derived from base types (like NUMBER) with additional constraints.

- **POSITIVE:** Restricts a number to positive values (e.g., `age POSITIVE`)
- **NATURAL:** Restricts a number to non-negative values (e.g., `balance NATURAL`)

**2. Collections:** Group similar data items together.

- **VARRAY:** Ordered collection with a fixed size at declaration.
- **TABLE:** Associative collection with key-value pairs.

**3. Large Objects (LOBs):** Store large data like text, images, or audio.

- **BLOB:** Binary Large Object.
- **CLOB:** Character Large Object.

**4. Record Types:** Group related variables of different data types into a single unit.

**5. Reference Types:** Store references to existing data structures. (e.g., `emp_cursor REF CURSOR`; (reference to a cursor))

## Variables

- Variables act as temporary storage locations for data during program execution.
- They allow you to manipulate data within your code.

**Declaring Variables:** Variables need to be declared using the **DECLARE** keyword. Declaration specifies the variable name, data type, and optionally an initial value.

## Constants

Constants are fixed values that cannot be changed during program execution.

**Declaring Constants:** using the **CONSTANT** keyword.

```
CONSTANT constant_name data_type := value [NOT NULL];
```

## Attributes in PL/SQL

**1. Cursor Attributes:** Provide information on a cursor's status after execution.

- **%ISOPEN:** Checks if the cursor is open.
- **%FOUND:** Indicates if a FETCH retrieved a row (TRUE) or not (FALSE).
- **%NOTFOUND:** Opposite of %FOUND (TRUE if no row fetched, FALSE otherwise).
- **%ROWCOUNT:** Returns the number of rows fetched using the cursor.

**2. %TYPE Attribute:** Helps in declaring variables based on existing data types. Ensures type compatibility between the PL/SQL variable and the referenced source.

# Control Structures

**1. Conditional Control:** Determines which block of code to execute based on a boolean condition

- **IF-THEN-ELSE Statement:** Executes one of two code blocks based on a condition.
- **CASE Statement:** Evaluates an expression against multiple conditions and executes corresponding code blocks.

**2. Iterative Control:** Executes a block of code repeatedly until a specific condition is met.

- **LOOP Statements:** Repeatedly execute a block of code until a condition is satisfied.
- **FOR Loop:** Iterates a specific number of times based on a counter variable.
- **WHILE Loop:** Continues execution as long as a condition remains TRUE.

**3. Sequential Control:** Executes statements in the order they appear in your code

## Cursors in PL/SQL

- A cursor acts as a pointer to a result set retrieved from a SELECT statement.
- It allows you to process data one row at a time, unlike fetching the entire set upfront.

**Types of Cursors:**

**1. Implicit Cursors:** Created automatically by PL/SQL for any SELECT statement.

**2. Explicit Cursors:** Defined by the programmer for granular control over data retrieval.

**Steps to Use Explicit Cursors:** Declare the Cursor > Open the Cursor > Fetch Data > Close Cursor.

**Benefits of Explicit Cursors:** Enhanced control over data processing. Memory efficiency when dealing with large datasets.

## Exception Handling

**Exceptions:** Errors that disrupt normal program flow during execution.

**Two Types:**

- **System-defined Exceptions:** Predefined errors by PL/SQL.
- **User-defined Exceptions:** Programmer-created exceptions for specific error conditions.

**Key Concepts:**

- **RAISE Statement:** Explicitly raises an exception to signal an error.
- **Exception Block:** Captures and handles raised exceptions.

## Triggers

Triggers are PL/SQL blocks stored in the database and can be invoked repeatedly. They are automatically executed upon specific events.

**Event-driven execution:** Triggers fire in response to various database events, including:

- **DML (Data Manipulation Language) events:** INSERT, UPDATE, DELETE (DML)
- **DDL (Data Definition Language) events:** CREATE, ALTER, DROP (DDL)
- **Database events:** Errors, logins/logouts, startup/shutdown

### Trigger Timing:

- **BEFORE:** Runs code to validate data or set values **before** changes occur.
- **AFTER:** Executes code to audit changes or react to them **after** they happen.
- **INSTEAD OF:** Less common, replaces the standard DML operation with custom logic.

### Trigger Benefits:

- **Ensure data quality:** Validate data and prevent invalid entries.
- **Automate tasks:** Reduce manual work by automating.
- **Enhance security:** Restrict data access and implement additional security measures.
- **Maintain data consistency:** Keep data synchronized across multiple tables.

## Procedures

Procedures are reusable blocks of PL/SQL code that perform specific tasks. They promote modular design by encapsulating functionality.

**Structure:** A procedure consists of a header (specification) and a body:

- **Header:** Defines the procedure name, optional parameters (IN for input, OUT for output), and return type (if applicable).
- **Body:** Contains the executable code (declarations, statements, exception handling).

### Benefits:

- **Reusability:** Procedures can be called from other PL/SQL code, reducing redundancy.
- **Modularity:** Break down complex tasks into smaller, manageable units.
- **Maintainability:** Easier to modify and debug code within procedures.
- **Performance:** Can improve performance by reducing network traffic.

## Packages

- A way to organize and group related PL/SQL code elements. Think of them as reusable modules containing procedures, functions, variables, cursors, and more.

### Benefits:

- **Modularity:** Break down large codebases into smaller, manageable units.
- **Reusability:** Share common code across different parts of your application.
- **Maintainability:** Easier to understand, modify, and debug code.
- **Encapsulation:** Control access to code elements (public vs. private).

### Package Specification:

- Declares the public interface (available elements).
- Includes declarations for procedures, functions, variables, cursors, etc.

### Package Body (Optional):

- Contains the implementation (code) for public elements.
- Also includes private declarations for helper functions or variables.

### Using Packages:

- Reference public elements using the `package_name.element_name` syntax.